

Contents

1	Graph Algorithms	2
1.1	Lecture 1: Single-Source Shortest Paths	2

1 Graph Algorithms

1.1 Lecture 1: Single-Source Shortest Paths

We approach the SSSP problem with possibly negative edge weights. As input, we get the directed graph $G = (V, E, w)$ with weight function $w : E \rightarrow \mathbb{R}$ and start vertex $s \in V$. We also take for a set of edges S , $w(S) = \sum_{s \in S} w(s)$. The algorithm should output the lengths of the shortest paths from s to any other $v \in V$ (or the shortest path tree). We know a few algorithms for this.

- Dijkstra (with a Fibonacci Heap): $O(m + n \log n)$, which only works if $w \geq 0$.
- Bellman-Ford: $O(mn)$.
- Bernstein-Nanongkai-Wulff-Nilsen: $\tilde{O}(m \log W)$, where $|w| \leq W$.

The $\tilde{O}(f)$ means $f \cdot \text{Polylog}(f)$. We discuss the third one today.

We define a price function as a function $\phi : V \rightarrow \mathbb{R}$ and the associated price-reduced weight function as:

$$w_\phi((u, v)) = w((u, v)) + \phi(u) - \phi(v)$$

Note the following observations. For any path $P = (v_0, \dots, v_r)$, $w_\phi(P) = w(P) + \phi(v_0) - \phi(v_r)$. As a corollary, for all cycles C , $w(C) = w_\phi(C)$. This implies the shortest path in $G_\phi = (V, E, w_\phi)$ is the same as in G . Our goal is thus to find a ϕ such that for all edges, $w_\phi \geq 0$, then reduce to Dijkstra.

Theorem 1.1

A ϕ satisfying $w_\phi \geq 0$ exists if and only if there exist no negative cycles in the original graph G .

Proof

Clearly if such a ϕ exists, $w_\phi(C) \geq 0$ for any cycle which is the same as $w(C)$ by our observations above. For the other direction, let $\phi(v) = d(s, v)$, where $d(s, v)$ is the length of the shortest path from s to v . This means taking some neighbor u of v ,

$$\phi(v) = d(s, v) \leq d(s, u) + w(u, v)$$

This means that $w_\phi(u, v) = w(u, v) + \phi(u) - \phi(v) \geq w(u, v) + d(s, u) - d(s, u) - w(u, v) = 0$.

The first person to use this technique was [G '75]. He showed the following result.

Theorem 1.2

If you have an algorithm that finds a good ϕ in time $T(m)$ for special case when $w(e) \geq -1$, then you can solve the general case of $|w| \leq W$ in time $O(T(m) \log W)$.

Proof

Round W up to the nearest power of 2, $W = 2^k$. We proceed by induction on k . If $k = 0$, all the edge weights are already at least -1 , so we're finished.

For an inductive case, we use the following algorithm, using solve as our subroutine:

```

function A( $G = (V, E, w)$ ,  $k$ )
  if  $k = 0$  then
    return Solve( $G$ )
  else
     $\hat{w}(e) \leftarrow \left\lceil \frac{w(e)}{2} \right\rceil$ 
     $\hat{\phi} \leftarrow A(V, E, \hat{w}, k - 1)$ 
     $\phi \leftarrow 2\hat{\phi}$ 
    return Solve( $G_\phi$ )

```

To prove correctness, we need to show that G_ϕ is solvable. Take some edge $e = (u, v)$.

$$\begin{aligned}
 w(e) &\geq 2 \left\lceil \frac{w(e)}{2} \right\rceil - 1 \\
 &= 2\hat{w}(e) - 1 \\
 w_\phi(e) &= w(e) + \phi(u) - \phi(v) \\
 &\geq 2\hat{w}(e) - 1 + 2\hat{\phi}(u) - 2\hat{\phi}(v) \\
 &= 2\hat{w}_\phi(e) - 1 \\
 &\geq -1
 \end{aligned}$$

Since the weights are nonnegative with the good $\hat{\phi}$.

Now, this is the novel portion. We will focus on G_s which is G with a dummy vertex s which has weight-0 to everyone else. We will focus on finding a ϕ for G_s (which we will label as G for brevity).

The algorithm has the following ingredients:

1. The subroutine Low Diameter Decomposition. $LDD(G, D)$ takes in a graph G where $w \geq 0$, and outputs $E^{rem} \subseteq E$ such that every strongly connected component of $G \setminus E^{rem}$ has weak diameter at most D . The weak diameter is $\max_{u, v \in \text{same SCC}} d_G(u, v)$. Furthermore, $\mathbb{P}[e \in E^{rem}] \leq O\left(\frac{w(e) \log^2 n}{D} + n^{-10}\right)$. This is a fast algorithm, with running time $\tilde{O}(m)$.
2. The subroutine Fix DAG Edges. This finds ϕ that makes $w_\phi \geq 0$ when G is a DAG. Using an SCC graph, this will mean $w_\phi(e) \geq 0$ for e going across SCCs. To implement this in linear time, just find the distance from a source with dynamic programming and just set the price function to be that.
3. The subroutine Elim Neg. This finds a ϕ in time $O(\log n \cdot \sum_v (1 + \eta_G(v)))$ where $\eta_G(v)$ is the number of negative edges on the shortest path to v . This algorithm assumes all in-degrees and out-degrees are $O(1)$.
4. The subroutine Scale Down. It takes in two numbers Δ and B . Assumes that $\eta(G) = \max_{v \in V} \eta_G(v) \leq \Delta$ and assumes all edges $w(e) \geq -2B$. It outputs ϕ such that $w_\phi \geq -B$.

Furthermore, in the original graph we can assume WLOG, all degrees are $O(1)$, so we can actually use the third condition. We do this with graph blow-up on each vertex v . Suppose v has x in-degree and y out-degree. We can make a $x + y$ -cycle with all edge weights 0. Each vertex has one of v 's original edges, either an incoming or outgoing one. This does not change shortest paths. All this does is blow up the number of vertices to $O(m)$ at most, which doesn't affect the runtime given.

Let's put them all together.

Algorithm 1.1

function MAIN($G = (V, E, w)$)

$B \leftarrow 2n$ (rounded up to nearest power of 2)

$\bar{w} \leftarrow Bw$

$\phi_0 \leftarrow 0$

for $i = 1$ to $\log_2 B$ **do**

$\psi_i \leftarrow \text{ScaleDown}((V, E, \bar{w}_{\phi_{i-1}}), \Delta = n, \frac{B}{2^i})$

$\phi_i \leftarrow \phi_{i-1} + \psi_i$

$w^* \leftarrow \frac{\bar{w}_{\phi_{\log B}}}{B} + \frac{1}{B}$